

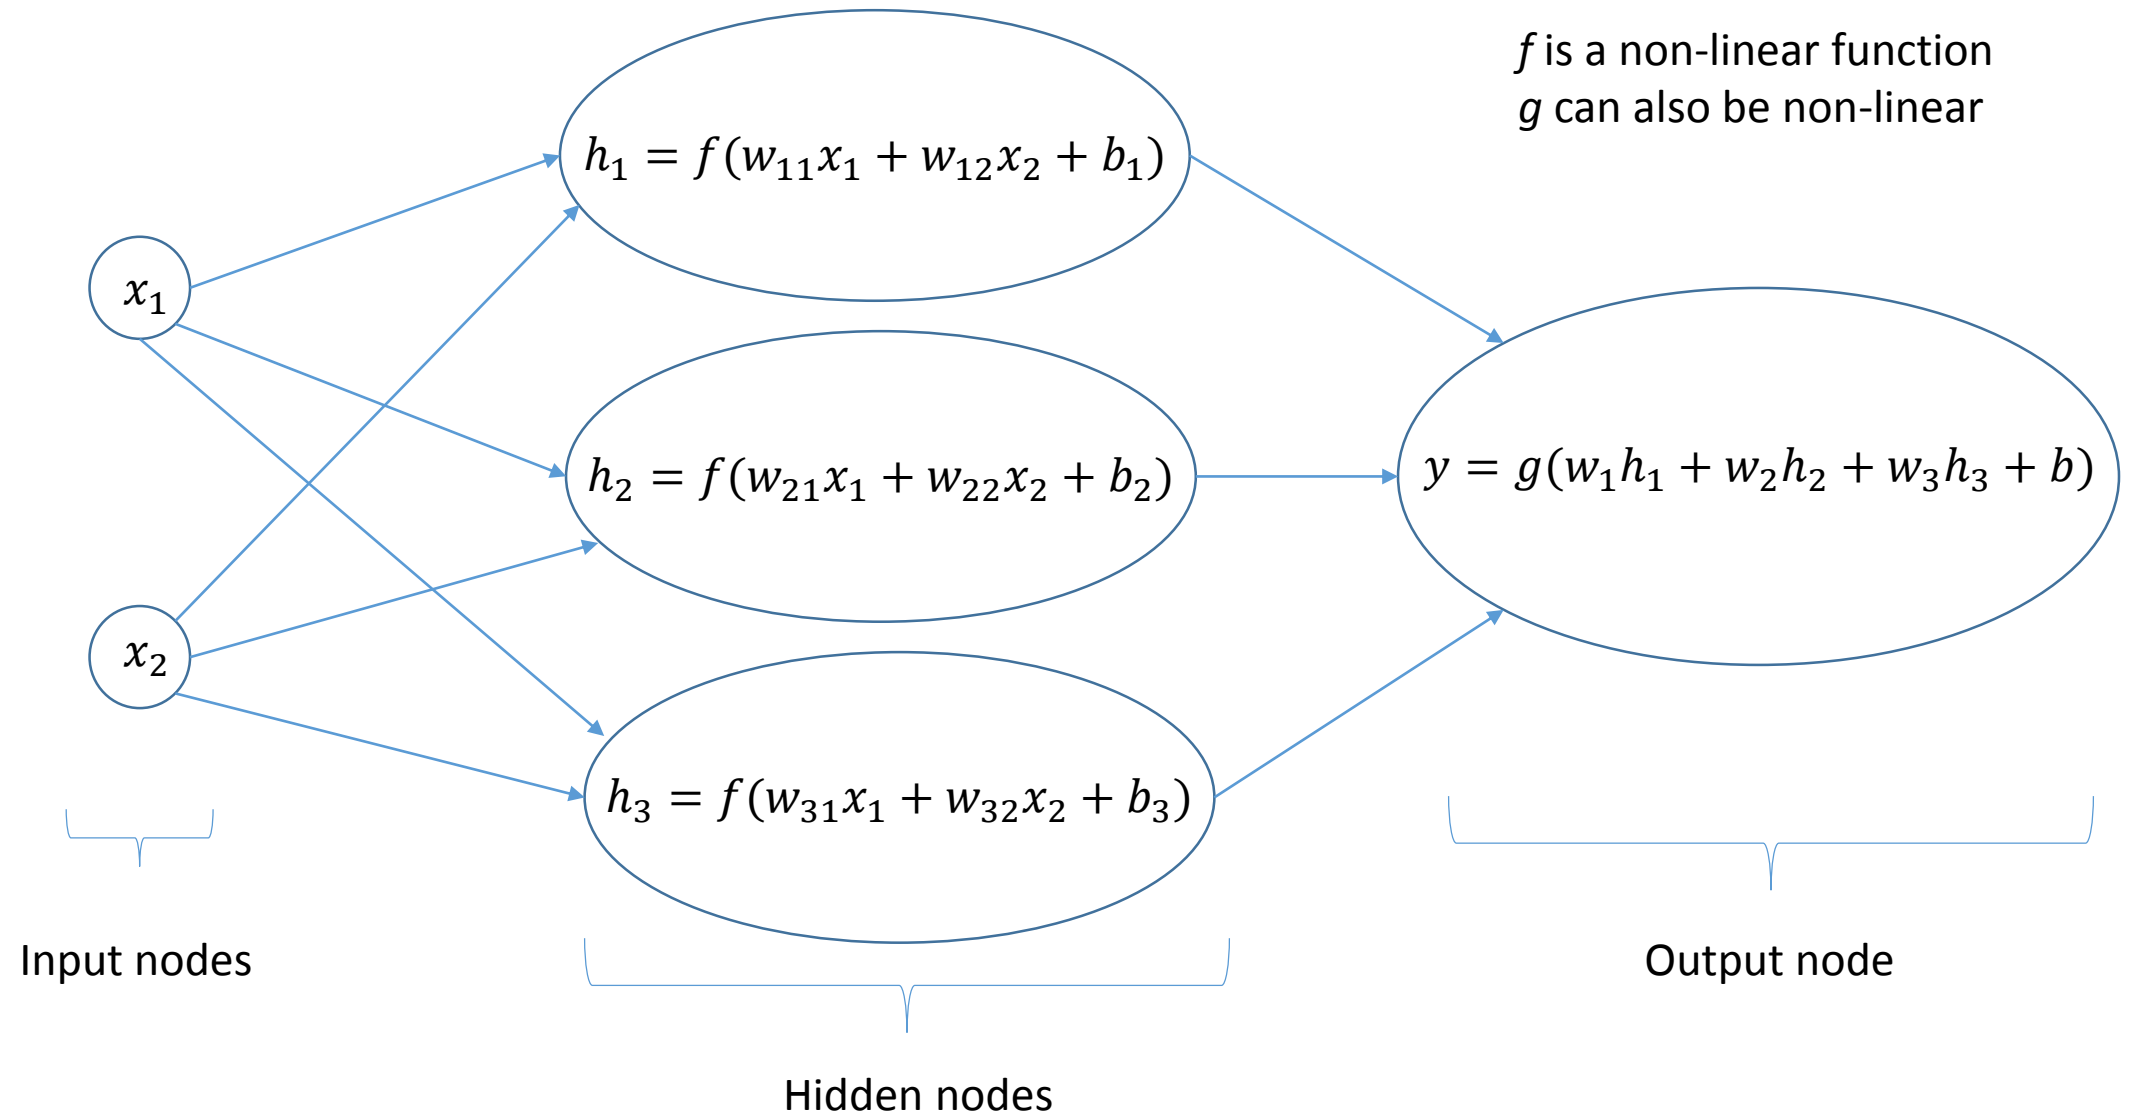
Introduction to Neural Networks

Computing Science
University of Alberta
Nilanjan Ray

Agenda

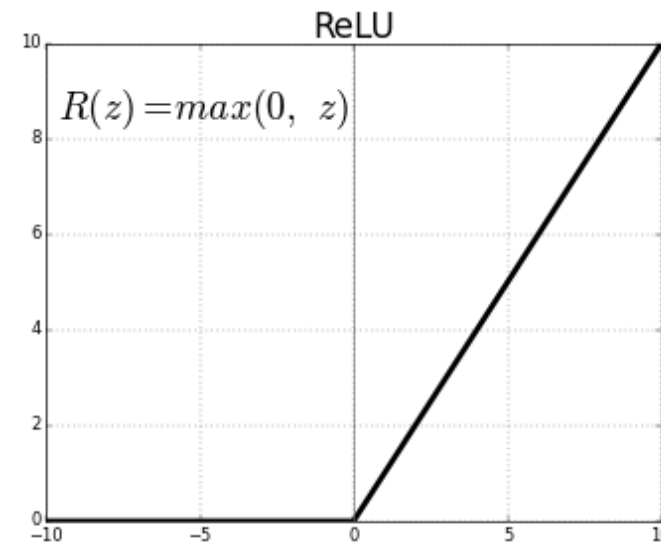

- What is a Neural Net?
 - Neural net as a computational graph
- Approximating “XOR” function with neural net
- Understanding backpropagation
- Universal function approximation by a neural net

Feed forward neural network




Feed forward net: non-linear functions

- Non-linear functions at hidden nodes are known as “activation function”
 - Sigmoid, ReLU, ELU,



Why activation functions are non-linear?

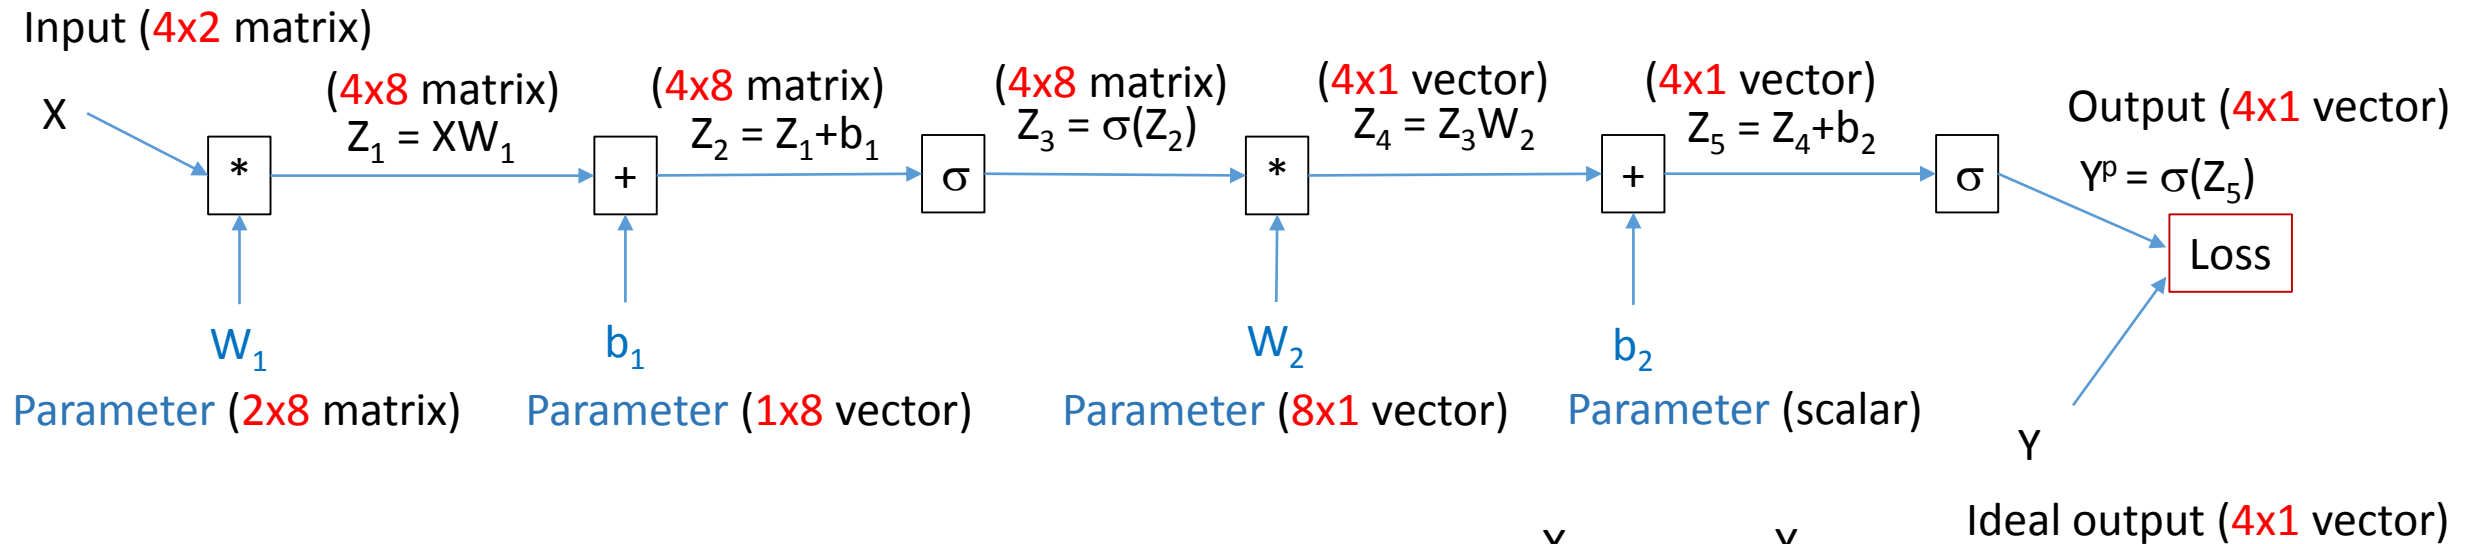
Feedforward net in general: Directed acyclic graph



What's the big deal about neural net?

- Mathematically very rich: it can approximate any function
- It is biologically inspired: (loosely) resembles brain connections
- Computationally:
 - Simple: matrix-vector multiplication and point-wise non-linear function
 - Highly parallelizable: cuBLAS, GEMM, Batched GEMM!
- Excellent **empirical** results on “generalization capability” over variety of applications!

Neural network as a computational graph



$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

Sigmoid function;
applied **pointwise**
to a vector or
matrix input


This network is
trying to learn
XOR function

X		Y
1	0	1
0	0	0
0	1	1
1	1	0

How does PyTorch optimize parameters?

- By **gradient descent** PyTorch adjusts network parameters to reduce the value of the loss function.
- But how?
 - Answer: **Backpropagation**
- **Let us learn how to do backpropagation on a computational graph!**

Chain rule of derivative for a computational node




If X , Z , W are all scalars, then usual chain rule of derivative applies:

$$\frac{\partial(\text{Loss})}{\partial X} = \frac{\partial Z}{\partial X} \frac{\partial(\text{Loss})}{\partial Z}$$

$$\frac{\partial(\text{Loss})}{\partial W} = \frac{\partial Z}{\partial W} \frac{\partial(\text{Loss})}{\partial Z}$$

Chain rule of derivative...




If X , Z , W are **matrices or vectors**, then :

$$\nabla_X(\text{Loss}) = \left(\frac{\partial Z}{\partial X} \right) * \nabla_Z(\text{Loss})$$

$$\nabla_W(\text{Loss}) = \left(\frac{\partial Z}{\partial W} \right) * \nabla_Z(\text{Loss})$$

“*” refers to matrix vector multiplication

Example 1




Chain rules:

$$\nabla_X(\text{Loss}) = W^T \frac{\partial(\text{Loss})}{\partial Z} = [w_1 \quad w_2] \frac{\partial(\text{Loss})}{\partial Z}$$

$$\nabla_W(\text{Loss}) = X^T \frac{\partial(\text{Loss})}{\partial Z} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \frac{\partial(\text{Loss})}{\partial Z}$$

Why?

Example 2




$$\nabla_X(\text{Loss}) = \nabla_Z(\text{Loss})W^T$$

Chain rules:

$$\nabla_W(\text{Loss}) = X^T \nabla_Z(\text{Loss})$$

Why?

Backprop derivation



Backprop derivation...

The diagram illustrates the backprop derivation through a series of steps, with annotations explaining each part:

- Step 1:** $\delta X_i = \sum_k \underbrace{\frac{\partial Z_k}{\partial X_i} \frac{\partial(\text{Loss})}{\partial Z_k}}_{\text{Chain rule of derivative}} = \sum_k \underbrace{\frac{\partial}{\partial X_i} \left[\sum_j X_j W_{jk} \right]}_{\text{Substitute } Z_k} \delta Z_k = \sum_k W_{ik} \delta Z_k$
 - δX_i is the i^{th} component of δX vector.
 - The first bracket indicates the chain rule of derivative.
 - The second bracket indicates substituting Z_k .
 - δZ_k is the k^{th} component of δZ vector.
- Step 2:** A large blue arrow points from the scalar equation to the matrix equation.
- Step 3:** $\delta X = \delta Z W^T$ (written in red).
- Step 4:** An annotation "Because," points to the matrix equation, with the supporting derivation: $\frac{\partial}{\partial X_i} \left[\sum_j X_j W_{jk} \right] = W_{ik}$.
- Step 5:** An annotation "Writing in matrix-vector multiplication form" points to the final matrix equation.

Backprop derivation...

$$\delta W_{ij} = \sum_k \underbrace{\frac{\partial Z_k}{\partial W_{ij}} \frac{\partial(\text{Loss})}{\partial Z_k}}_{\text{Chain rule of derivative}} = \sum_k \frac{\partial}{\partial W_{ij}} \underbrace{\left[\sum_m X_m W_{mk} \right]}_{\text{Substitute } Z_k} \delta Z_k = X_i \delta Z_j$$

$(i,j)^{\text{th}}$ component of δW matrix

Chain rule of derivative

Substitute Z_k

Because,

$$\frac{\partial}{\partial W_{ij}} \left[\sum_m X_m W_{mk} \right] = \begin{cases} X_i, & \text{if } i = m \text{ and } j = k, \\ 0, & \text{otherwise.} \end{cases}$$


k^{th} component of δZ vector

$\delta W = X^T \delta Z$

Writing in matrix-vector multiplication form

Backprop derivation...

“Broadcast” addition:



$$\delta X_{i,j} = \sum_k \sum_l \frac{\partial Z_{k,l}}{\partial X_{i,j}} \delta Z_{k,l} = \sum_k \sum_l \frac{\partial}{\partial X_{i,j}} [X_{k,l} + b_k] \delta Z_{k,l} = \delta Z_{i,j} \quad \Rightarrow \quad \delta X = \delta Z$$

Chain rule


Substitute $Z_{k,l}$

Because,

$$\frac{\partial}{\partial X_{i,j}} [X_{k,l} + b_k] = \begin{cases} 1, & \text{if } i = k \text{ and } j = l, \\ 0, & \text{otherwise.} \end{cases}$$

Backprop derivation for broadcast addition

“Broadcast” addition:



$$\delta b_i = \sum_k \sum_l \underbrace{\frac{\partial Z_{k,l}}{\partial b_i}}_{\text{Chain rule}} \delta Z_{k,l} = \sum_k \sum_l \frac{\partial}{\partial b_i} [X_{k,l} + b_l] \delta Z_{k,l} = \sum_k \delta Z_{k,i}$$

Substitute $Z_{k,l}$

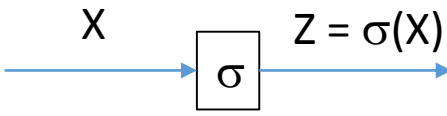
Because,

$$\frac{\partial}{\partial b_i} [X_{k,l} + b_l] = \begin{cases} 1, & \text{if } i = l, \\ 0, & \text{otherwise.} \end{cases}$$

$$\delta b = \sum_k \delta Z_{k,:}$$

Backprop derivation for activation function

Non-linear function:
(applied **pointwise**)



Using chain rule: $\delta X_{i,j} = \frac{dZ_{i,j}}{dX_{i,j}} \delta Z_{i,j} = \frac{d\sigma(X_{i,j})}{dX_{i,j}} \delta Z_{i,j}$

If the non-linear
function is sigmoid,

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

$$\frac{d\sigma}{da} = \frac{\exp(-a)}{(1 + \exp(-a))^2} = \frac{1}{1 + \exp(-a)} \left(1 - \frac{1}{1 + \exp(-a)} \right) = \sigma(a)(1 - \sigma(a))$$

$$\delta X_{i,j} = \sigma(X_{i,j})(1 - \sigma(X_{i,j}))\delta Z_{i,j}$$

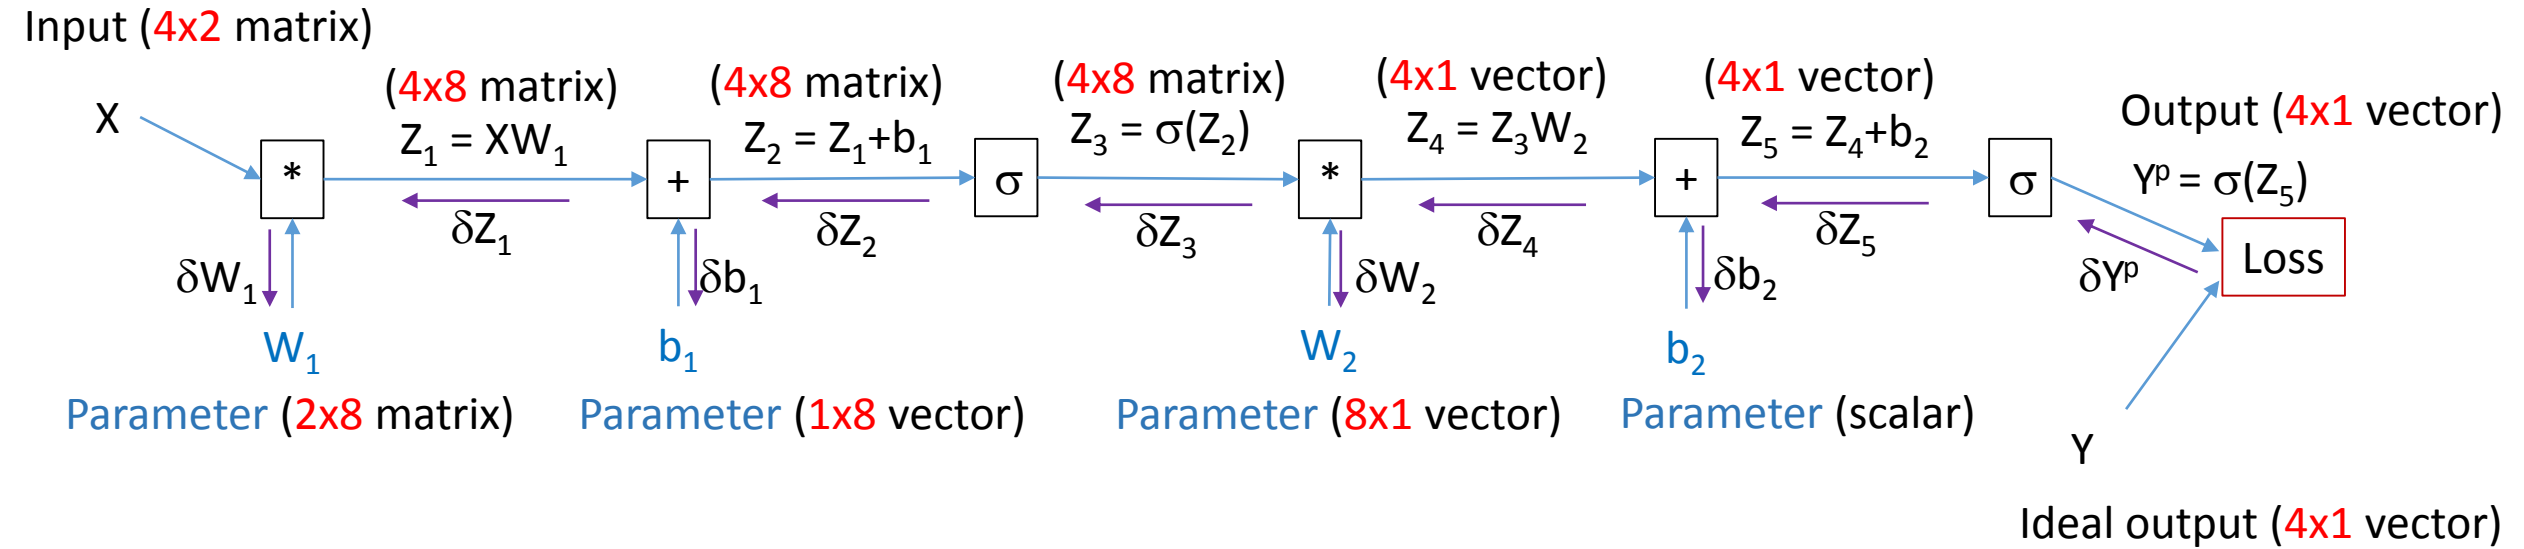
Backprop derivation for loss function

Euclidean loss function: $Loss(Y^p, Y) = \frac{1}{2} \|Y^p - Y\|^2 = \frac{1}{2} \sum_i (Y_i^p - Y_i)^2$

i^{th} component of δY^p vector: $\delta Y_i^p = \frac{\partial}{\partial Y_i^p} Loss(Y^p, Y) = \frac{\partial}{\partial Y_i^p} \frac{1}{2} \sum_i (Y_i^p - Y_i)^2 = Y_i^p - Y_i$

Using vector notation: $\delta Y^p = Y^p - Y$

Apply chain rule to XOR neural network



Chain rule of derivatives:

$$\delta Y^p = Y^p - Y$$

$$\delta Z_5 = \sigma(Z_5)(1 - \sigma(Z_5))\delta Y^p$$

$$\delta Z_4 = \delta Z_5$$

$$\delta Z_3 = \delta Z_4 W_2^T$$

$$\delta Z_2 = \sigma(Z_2)(1 - \sigma(Z_2))\delta Z_3$$

$$\delta Z_1 = \delta Z_2$$

Gradient of "Loss" with respect to input signals



Propagates backward

$$\delta W_2 = Z_3^T \delta Z_4$$

$$\delta b_2 = \sum_k (\delta Z_5)_k$$

$$\delta W_1 = X^T \delta Z_1$$

$$\delta b_1 = \sum_k (\delta Z_2)_{k,:}$$

Gradient of "Loss" with respect to parameters

New notation:
 $\delta S \equiv \nabla_S(\text{Loss})$


Backprop to train a neural net

Initialize all parameters of the neural network

Initialize learning rate variable lr

Iterate:

If loading the whole training data, do it **once** outside the “Iterate” loop, to be efficient



(Load Data): Get training data batch

(Forward pass): Compute Z_1, Z_2, \dots, Y^p

(Backward pass): Compute gradients $\delta Y^p, \delta Z_5, \dots, \delta Z_1, \delta W_2, \delta W_1, \delta b_2, \delta b_1$

(Gradient descent to update parameters): $W_2 \leftarrow W_2 - lr * \delta W_2, \quad b_2 \leftarrow b_2 - lr * \delta b_2, \dots,$

(Diagnostics): Compute “Loss” from time to time to check if it is decreasing

“Learn_XOR_manualBP.ipnyb” implements this learning algorithm

PyTorch magic!

- Fortunately, PyTorch can automatically compute derivatives by chain rules!
- Also, it has several optimizers that can use these derivatives in the gradient descent optimization method.
- Look at “Learn_XOR.ipnyb” and “Learn_XOR_with_LBFGS.ipnyb”

Universal function approximation

- A neural network with a single hidden layer can approximate “any” function!
 - Wikipedia has a clear statement:
https://en.wikipedia.org/wiki/Universal_approximation_theorem
- A non-technical explanation of universality theorem:
 - <http://neuralnetworksanddeeplearning.com/chap4.html>

Why do we then even need multiple layers and why even deep nets?